

THE AI MASTERCLASS

Exclusive Resources PDF

March 2025




Index

Index	2
AI Presentation & Visual Design Tools	3
AI Website Building Tools	5
AI Image Generation Tools	8
AI Video Generation Tools	15
AI Audio Generation Tools	22
AI Coding Assistant	25
Free Tools List	27

AI Presentation & Visual Design Tools

These tools help create presentations or visual content quickly using AI.

Napkin AI (Presentation/Diagram Generation)



Description: Napkin AI turns text into visuals like [diagrams](#) or [infographics](#), aiding business storytelling and presentations.

It's great for quickly visualising concepts from plain text.

- **Pros:** Very easy to use; offers great customisation of visuals; extremely quick to generate results. Integrates as a **document editor** for visuals, allowing [quick](#) conversion of ideas to charts or graphics. Free [beta](#) access (no cost during beta period) encourages exploration.
- **Cons:** Currently **limited templates** available, so users might find only certain types of visuals. Struggles with unstructured input (it performs best with clear textual descriptions). As a newer tool, some features (like offline use or advanced editing) may be lacking.
- **Pricing:** **Free during beta** (as of 2025) – Napkin offers a free Starter plan (up to 3 visuals, limited styles) and the full [Professional](#) features are also free while in beta. Eventually, premium plans are expected to start around ~\$8/month per user, but early users will get [special](#) deals.
- **Use Cases:** Creating quick infographics, flowcharts, diagrams; educators turning lesson text into visuals; business teams visualising processes or ideas without a graphic designer. Ideal for anyone needing a quick visual to accompany text content.
- **Official Site:** napkin.ai (Sign up for free beta)

Gamma (Presentation & Document Creator)



Description: [Gamma](#) is an AI-powered tool for [generating](#) **slide decks, documents, and web pages** from outlines or prompts. It can produce a formatted presentation with AI-generated text and images, and allows easy refinement of any slide.

- **Pros:** Extremely **user-friendly** interface for creating decks—just write a prompt and Gamma generates a polished [presentation](#). It has many templates and design options and supports interactive elements. Flexible output (presentations, PDFs, websites) and [real-time collaboration](#) features are available. The AI assistant helps refine content and layout on the fly.
- **Cons:** The **free plan is limited** – new users get [400](#) AI credits one-time and up to 10 slides per creation, which do not renew, meaning heavy users must upgrade. Not all advanced image generation models are available on free tier (basic models only). Also, presentations [cannot](#) be edited offline (internet required).
- **Pricing:** [Free plan](#) available (no cost, includes 400 AI credits and basic features). Paid plans include **Plus (\$10/month)** and **Pro (\$20/month)**, with discounts for annual billing (e.g. ~\$8/mo and ~\$15/mo respectively when paid yearly). Plus offers unlimited AI generation and custom branding, Pro adds more advanced image models and up to 60 cards (slides) per deck. *(Note: 400 free credits ≈ 20,000 token input and basic image generation).*
- **Use Cases:** Quickly drafting **pitch decks, project proposals, reports** with visual appeal. Non-designers can create slides by simply describing content. Also useful for building simple one-page websites or documents to share information interactively. Great for startups, teachers, or anyone needing a fast presentation.
- **Official Site:** [gamma.app](#) (offers free signup and examples).

Comparison – AI Presentation Tools:

Tool	Ease of Use	Customization Options	Free Tier Limits	Notable Features	Best Use Case
Napkin AI	Very high (simple text-to-visual)	Basic editing (colours, 3 preset styles on free); more coming post-beta	Free during beta (3 visuals limit)	Instantly creates diagrams/infographics from text	Visualizing ideas or data quickly in docs/presentations
Gamma	Very high (prompt-based deck creation)	Good design templates; advanced image generation on paid plans	Free forever tier (400 AI credits, ~10 slides)	Exports to web, PDF, PPT; AI refines content and layout	Generating polished slide decks or docs from an outline fast

AI Website Building Tools

These AI tools assist in **building websites** or web content with minimal coding.

Framer (AI Website Designer & Builder)



Description: Framer is a web design tool that now offers an **AI [website builder](#)**. By inputting a description (or using AI to generate text/images), Framer can create a multi-page, responsive website in seconds. It combines design and deployment – users can further tweak the design in Framer’s editor and publish with built-in hosting.

- **Pros: Easy and intuitive** visual editor – users find it more straightforward than traditional site builders like WordPress or Wix. Combines design and building in one (similar to Figma’s flexibility but outputs a live site). Supports complex interactions, animations, and produces fast, SEO-optimised sites out of the box. Great integration of AI [copywriter](#) for text content in your [brand](#) voice.
- **Cons:** Not suitable for web applications or sites needing complex backend logic (it’s primarily for static or front-end oriented sites). Uses a **subscription model** for custom domains and advanced features. Also, while easier than code, the design interface is similar to design tools (Photoshop/Figma), so **absolute beginners** might face a learning curve understanding design concepts.
- **Pricing: Free plan** ("[Free Forever](#)") allows one Framer subdomain site with up to [1,000](#) monthly visitors (great for simple personal sites). Paid plans are per site: **Mini** – ~\$5/month (annual) for custom domain, limited pages; **Basic** – ~[\\$15/month](#) for up to 150 pages and more CMS items ; **Pro** – ~\$25/month for larger sites (300 pages, more traffic). All paid tiers remove Framer branding and increase visitor limits. (*Framer’s design tool usage itself is free; pricing applies when hosting/publishing sites*).
- **Use Cases:** Designers and marketers building **landing pages, portfolios, marketing sites** quickly. Small businesses that need a sleek website without hiring developers. Also useful for prototyping website ideas – you can have a working site from an AI prompt, then refine design for production.
- **Official Site:** [framer.com](#) (AI Site Builder is accessible on the website)

10Web AI Builder (WordPress Website Generator)




Description: 10Web is a platform offering AI-assisted WordPress site building and hosting. It can **recreate**

any website from a URL or generate a new site based on user prompts, then let you edit with a familiar WP/Elementor interface. It handles the full stack: AI content generation, design, and managed hosting.

- **Pros:** Seamless **WordPress integration** – sites are built on WordPress with Elementor, so you have a standard CMS after AI generates the initial design. Generally easier and faster than building a WP site from scratch; AI generates text and images that are surprisingly relevant. Great for quickly cloning a site's design: the **AI Recreation Assistant** can import an existing site's layout and style to jumpstart development.
- **Cons:** The AI generation can produce **technical errors** or imperfect layouts, so you might need to fine-tune (the tool is improving over time). Customisation options are somewhat limited to what Elementor offers – for highly bespoke designs or features, you might hit constraints. It also uses a monthly subscription model (includes hosting), which may not suit those who prefer one-time payments for a site.
- **Pricing:** Offers a **7-day free trial** (full access).
After that, subscription plans bundle hosting + builder: **AI Starter** at ~\$20/month (or ~\$10/mo annual) for 1 site with basic resources; **AI Premium** ~\$30/mo (or ~\$15/mo annual) for more robust sites ; higher plans like **AI Ultimate** for \$45/mo (more traffic) , and agency plans for multiple sites (e.g., 10Web Personal \$14/mo for 1 site, Premium \$35 for 3 sites, etc). *All plans include the AI builder, hosting, and site optimisations.*
- **Use Cases:** Perfect for quickly building a **WordPress website** for a small business, blog, or eCommerce (10Web has eCommerce plans for online stores as well). Agencies can use it to spin up client sites fast, then refine content. Also useful to migrate an existing site into a modern WP design using AI.
- **Official Site:** 10web.io (AI Website Builder information and trial available).

ChatGPT (for Website Code Generation)



Description: *ChatGPT by OpenAI* – a powerful language model – isn't a dedicated site builder, but it was highlighted for its ability to **assist in website creation** by generating code and content. Through natural language prompts, ChatGPT (especially with GPT-4 and the new Vision feature) can produce HTML/CSS/JS and even interpret sketches of a layout to help build a basic website

- **Pros: Excellent brainstorming partner** – can suggest website ideas, features, and write boilerplate code or text content. With GPT-4's code capability, it can often convert written instructions into working website code (for example, generate a simple HTML/CSS page from a prompt). The Vision feature (if available) can interpret images or rough drawings of a site and output corresponding code structure. It's also readily available (free tier) and supports many frameworks via code plugins.
- **Cons:** Not a specialised tool for design – ChatGPT won't give you a polished visual layout out-of-the-box; you'll likely need to style the generated code. It may produce code that needs debugging or

refining, especially for complex layouts or interactive features. Also, it does not *deploy* a site (it just gives code), so you must have the ability to use the code it provides. In short, it's "helpful but not a one-click site builder."

- **Pricing: Free to use** for anyone on the web (uses the GPT-3.5 model for free). For the most advanced model and features (like GPT-4, GPT o1, o3 with vision, priority access), **ChatGPT Plus** costs **\$20 per month**. The Plus subscription includes GPT-4, faster responses, and plugin usage, which significantly boost its coding assistance abilities. *(No additional cost per use, except if using OpenAI API which is usage-based.)*
- **Use Cases:** Useful for developers or enthusiasts to **generate snippets of code**, fix HTML/CSS issues, or get quick examples of how to implement something on a webpage. For non-coders, it can create a starting point for a simple site ("a single-page personal website") which they can then hand off to a developer or refine themselves. Also great for generating content text for sites (product descriptions, blog sections, etc.).
- **Official Site:** chat.openai.com (ChatGPT interface – Plus subscription via OpenAI)

Comparison – AI Website Building Tools:

Tool	Platform Type	AI Capability	Pricing (starting)	Strengths	Limitations
Framer	All-in-one design & hosting	Generates multi-page responsive sites from prompt AI writes copy	Free (Framer subdomain); \$15/mo Basic custom domain	Polished designs, animations, fast hosting	No backend logic, learning curve for complex design
10Web AI	WordPress-based builder/ hosting	Recreates sites from URL or text; AI generates WP pages content/ images	7-day free; \$20/ mo (or \$10/mo annual) single site	WordPress CMS power + AI speed; content & SEO tools built-in	AI output may need manual subscription includes hosting (not standalone builder)
ChatGPT	AI chatbot (coding assistant)	Writes code or content from natural language; can interpret sketches (GPT-4 Vision)	Basic version free; \$20/mo for Plus subscription; \$200/mo for Pro subscription	WordPress CMS power + AI speed; content & SEO tools built-in	You'd need to have knowledge of coding

AI Image Generation Tools

These tools create images from text prompts (text-to-image). They differ in models, quality, and usage policies.


DALL-E 3 (OpenAI Image Generator)



Description: DALL-E 3 is OpenAI's latest text-to-image model, integrated into [ChatGPT](#). It excels at understanding nuanced prompts by leveraging [GPT-4](#) to refine prompt details. Users can converse with ChatGPT to generate and iteratively refine images. DALL-E 3 creates a wide variety of art styles and illustrations, though its default style leans slightly more **artistic/animated** than [photorealistic](#).

- **Pros: Extremely easy to use** – simply describe what you want in plain English (especially via ChatGPT's interface) and it will produce images. Great at understanding complex scene descriptions and helpful in guiding users to phrase prompts (the ChatGPT integration can suggest improvements). DALL-E 3's integration in ChatGPT Plus means you can have a conversational **prompting experience**, iterating on images with feedback. Also, usage rights are user-friendly – images you create are yours to use freely (including commercially).
- **Cons:** Struggles to generate **realistic human faces of actual people** – it avoids producing faces of real individuals due to ethical guardrails. Has **strict content censorship**: it refuses prompts about famous people, violence, guns, political figures, etc., which can limit certain creative scenarios. Some users find its outputs a bit stereotypical or lacking in imaginative creativity for very abstract prompts– it might default to common interpretations. Also, outside of ChatGPT, there's no standalone DALL-E 3 interface except through Bing.
- **Pricing:** Accessible **for free on Bing** (Microsoft's Bing Image Creator uses DALL-E 3 for no charge). Within ChatGPT, **DALL-E 3 is included in the \$20/month ChatGPT Plus** subscription at no extra cost. There is no separate pricing for DALL-E 3 itself (OpenAI's older DALL-E 2 had pay-per-image credits, but DALL-E 3 is mainly through the ChatGPT Plus or Bing channels). Essentially, Plus users get a generous amount of image generation included.
- **Use Cases:** Ideal for quickly generating **illustrations, concept art, product mockups, or storybook-like images** from simple prompts. Great for marketing material ideation, blog post illustrations, or visual brainstorming. Non-artists can get drawings for their ideas (e.g. a book cover concept). The strong prompt understanding makes it suitable for complex scenes or when you need the AI to fill in details from a broad concept.
- **Official Link:** Via [ChatGPT Plus](#) (openai.com) or [Bing Image Creator](#) (free access via Microsoft)

Leonardo AI



Description: Leonardo.ai is a versatile image generation platform originally geared towards creating game assets, now expanded to general image art generation. It offers multiple models/finetunes and a feature-rich web app for generating and editing AI images. Users get fine control over styles, model selection (e.g. [Leonardo's](#) own models, Stable Diffusion variants), and even basic animation features.

- **Pros: Many more options and model choices** than basic generators- you can pick different AI models depending on desired style (fantasy art, illustrations, etc.), adjust dimensions, and use advanced settings. It adheres to prompts well, often yielding **clear, detailed images** as intended . It's considered on par with Adobe Firefly in quality and often better at following prompts than Midjourney in [CNET's](#) tests. Another plus: **low censorship** - Leonardo has minimal content restrictions (within reason), allowing more freedom in what you can create.
- **Cons:** The abundance of options can be **overwhelming for beginners**. The interface, while powerful, has many toggles and settings - new users may find it complicated and there's a learning curve to mastering it. Also, certain advanced features (like high-res upscaling or editing) are not available on the free plan. Finally, while prompt adherence is good, achieving ultra-photorealism may require choosing specific models or fine-tuning prompts; out of the box, some outputs might still have an illustrative quality compared to Midjourney's realism.
- **Pricing: Comprehensive free plan** - when you sign up, you automatically get a free tier with **150 tokens every 24 hours**, which is enough for dozens of images daily. This free usage renews continuously, making Leonardo one of the most generous free offerings. Paid plans (as of 2025) start at **\$10/month** for a Standard plan (if billed monthly, \$8 if yearly), which gives 700 monthly credits *plus* 30 daily credits refill (so ample generations) and perks like no watermark and priority processing. Higher tiers (Pro, etc.) cost more and offer even more credits or unlimited generations. For API usage or enterprise, pricing varies, but for most individuals the free or \$10 plan is sufficient.
- **Use Cases:** Great for artists, game designers, or creative enthusiasts who want **fine control**. It's used to create **game sprites, concept art, characters, environments**, as well as general digital art. Because of prompt fidelity, it's useful when you have a very specific image in mind. Also, educators or content creators who need a steady stream of images (Leonardo's free credits refresh means you can generate daily content). Its combination of **image and even basic video (animation)** tools makes it a multi-purpose creative suite.
- **Official Site:** [leonardo.ai](#) (Free account available; web-based studio).

Midjourney



Description: Midjourney is one of the leading AI image generators known for its **high-quality, imaginative outputs**. It operates via Discord (bot commands) and produces images in various styles, often with stunning detail and artistry. Midjourney is particularly praised for photorealistic images and creative art, making it a benchmark in AI art.

- **Pros: Excellent at realistic photography** as well as abstract art. Many users find Midjourney v5/v6 outputs to be among the most lifelike (for portraits, landscapes) and artistically rich. It offers great options for modifying images: you can upscale, create variations, and use new features like "Zoom Out" and "Pan" to extend images, and even achieve character consistency across images. It's continuously evolving – new features have reduced issues (e.g., it's getting better at consistent hands/faces). The community and support (via Discord) provide lots of inspiration and prompt ideas.
- **Cons:** Many features are tied to the **Discord interface**, which can be confusing for beginners. Using Midjourney means learning Discord bot commands and managing a feed of images, which some may find unintuitive. There is **no free tier** (after trial) – originally Midjourney had a free trial, but it's often off or limited, so effectively it's paid-only for regular use. Also, while censorship is moderate (it allows more than DALL-E but still filters some content), it does have some banned prompts and will not produce certain NSFW or harmful imagery. Finally, because Midjourney's style is so high-quality, people can often recognise "a Midjourney look," and achieving very specific styles may require prompt engineering.
- **Pricing: No ongoing free plan.** Midjourney requires a subscription: **Basic plan at \$10/month** (200 image generations per month), **Standard at \$30/month** (approximately 15 hours of fast generation, which usually means ~900 images, plus unlimited in relaxed mode). There is also a **Pro plan \$60/mo** (more fast hours, stealth mode for private generations) and a **Mega plan \$120/mo** for enterprise-level use. Annual billing gives ~20% discount (e.g., Basic \$96/year). *All plans are via the Midjourney website/ Discord sign-up.*
- **Use Cases:** Favoured by designers, artists, and creators for **concept art, product design mockups, album covers, fantasy illustrations, portraits, landscapes** – whenever top-tier visual quality is needed. Midjourney is often used in advertising and creative industries to prototype visuals. It's also popular in social media and content creation for unique art. If you need an AI image that could pass as professionally illustrated or photographed, Midjourney is a go-to.
- **Official Site:** midjourney.com (invites to Discord for usage; subscription required).

Adobe Firefly



Description: Adobe Firefly is Adobe's generative AI engine integrated into Creative Cloud apps (Photoshop, Illustrator, etc.). It can generate images from text, but its standout feature is integration with Adobe's tools – e.g., **Generative Fill** in Photoshop allows users to extend or modify images using AI. Firefly is tailored for design workflows, offering art styles that blend well into creative projects.

- **Pros:** Works **within Adobe apps** like Photoshop, making it convenient for designers to use AI directly in their existing workflow. It has **great options for modifying images** – generative fill, object removal, and style adjustments are all at your fingertips when editing file. Firefly also supports text effects and vector recolouring in Illustrator. For those already in the Adobe ecosystem, it's seamless. It produces consistent results and is easy to use with straightforward prompts, especially for concept art or design elements. And importantly, content generated in Firefly (non-beta features) is safe for commercial use with Adobe's IP indemnity, which is attractive to enterprise users.
- **Cons: Limited free usage** – free users get only a small number of generative credits (around 25 per month). After that, you need a subscription or it watermarks the images. In terms of output, while good, many find Firefly's image generation not as advanced as Midjourney or Leonardo for photorealism. As noted, it can be *"inferior in generation compared to Midjourney, Leonardo, etc."*, meaning if ultimate quality or variety is needed, other tools may outperform it. Firefly is also tightly integrated with Adobe's UI – it's not a standalone web service beyond beta testing, so you typically use it inside an Adobe app (Photoshop's web beta aside).
- **Pricing: Included with [Adobe Creative Cloud plans](#).** Users with a paid Adobe plan get a certain amount of **generative credits** each month (e.g., Photoshop subscribers might get some credits for Firefly features). Adobe also offers a **Firefly web** with a credit system: Free tier provides ~25 credits/month (with watermark). The **Firefly Premium** add-on was introduced at ~\$4.99/month for 100 credits (no watermark). Additionally, Creative Cloud enterprise plans include larger credit allotments. In summary, Firefly is generally accessed by having an Adobe subscription – if you already pay for Adobe, you likely have access; otherwise, a small monthly fee is needed for meaningful use.
- **Use Cases:** Perfect for graphic designers and marketers working in **Photoshop/Illustrator**: e.g., quickly expanding an image background, removing unwanted objects, generating variations of an image in a layout. Also useful for creating **custom stock images, marketing graphics, or concept art** that you will further refine in Adobe apps. If you "live" in Adobe tools, Firefly speeds up many tasks (like ideation and editing) without leaving your workspace.
- **Official Site:** [Adobe Firefly](#) (or use directly in Photoshop CC).

DreamStudio (Stable Diffusion by Stability AI)



Description: DreamStudio is the official web app for **Stable Diffusion**, the popular open-source image generation model by Stability AI. It provides an easy interface for generating images with Stable Diffusion and its latest versions (like SDXL). DreamStudio offers sliders for advanced parameters (steps, CFG, etc.) and supports using **negative prompts** to refine results. Essentially, it's the gateway to Stable Diffusion for non-coders.

- **Pros: Better customisation** than some closed models – users can tweak many generation settings (resolution, number of inference steps, guidance scale) to get the desired outputs. It also supports **negative prompts**, which let you specify what you *don't* want in the image (a powerful way to avoid certain elements). DreamStudio runs on Stability's cloud, so no local setup needed, but you can also take the same model and run it locally later if you want. New users get free credits (so you can try it out without payment). Because it's Stable Diffusion, there's a huge community making custom models and fine-tunes – DreamStudio often gives access to the latest official model releases.
- **Cons: Photorealism can be challenging for beginners** – out of the box, if you don't craft the prompt well, Stable Diffusion might give results that look more "AI-ish." The slide noted it's "bad at generating photos if you are a beginner", meaning it has a learning curve to get truly photo-quality results (unlike Midjourney which often does it automatically). There's a **massive learning curve** if you go beyond the app – i.e. running SD locally or training models requires technical skill. Also, the Stability team's default models have **high censorship** on certain content (though you could use community models to circumvent, the official platform will filter NSFW). In DreamStudio, you pay per generation, so heavy usage can consume credits quickly.
- **Pricing:** DreamStudio operates on a credit system. **New users get some free credits** (commonly reported as 25 free generations or 100 free credits). After that, you buy credits – the standard rate is about **\$10 for 1000 credits**, which yields roughly 500 images (since a default image might cost 2 credits for a lower resolution, more for higher res). For example, one source notes **25 free credits**, then \$10 for 1000 credits (enough for ~5,000 images if using the base model at low res). In practice, using SDXL at decent resolution might cost ~2 credits per image, so ~\$0.02 per image, which is quite cost-effective. No monthly subscription needed – it's pay-as-you-go. *(Alternatively, one can run Stable Diffusion locally for free, but with hardware costs.)*
- **Use Cases:** Good for AI enthusiasts or professionals who want **full control**. You can create any kind of art: from **concept art, anime, portraits, product designs, landscapes** – Stable Diffusion is very flexible especially if you find or train a model tuned to your domain. It's also used for AI **experimentation**: testing prompts, creating variations, or outputting a base image that you then edit manually. DreamStudio specifically is great for those who want to use Stable Diffusion without

installing anything, and for accessing the latest models (like SDXL) with a UI.

- **Official Site:** dreamstudio.ai (Stability AI's official interface for Stable Diffusion; requires login).

Comparison - AI Image Generators:

Tool	Model Type	Output Quality & Style	Free Usage	Notable Pros	Cons/ Cautions	Best For
DALL-E 3	Proprietary (OpenAI GPT-based)	High-quality; slightly <i>illustrative</i> look (great detail, sometimes "animated" feel)	Yes - Free via Bing; included in ChatGPT Plus	Easiest to use (ChatGPT integration); excellent prompt understanding; safe commercial use	Strong content filters; struggles with real faces; no standalone app (uses ChatGPT/ Bing)	Casual users; quick concept art; when needing AI to help craft the prompt
Leonardo AI	Proprietary + SD variants	High-quality; very controllable style (multiple models: from anime to realism)	Yes - 150 tokens/24h free (~15 images/day)	Many models & fine-tuning options; minimal censorship; prompt fidelity is strong	Can overwhelm new users (lots of settings); advanced features require paid plan	Power users; those wanting specific styles or assets (e.g. game art)
Midjourney	Proprietary (Discord-based)	Ultra high-quality; photorealistic or fantastical with artist-level detail	No (trial occasionally, but effectively paid)	Best photorealism and artistry in many cases; great image modification features (upscale, variations)	Paid only; requires Discord use; some learning curve with prompts; moderate filters (no extreme content)	Professionals needing top-tier visuals; creative projects where quality is paramount

Adobe Firefly	Proprietary (Adobe)	Good quality; excels in integrated editing (generative fill, etc.)	Limited - ~25 free credits/month (watermarked)	Integrates with Adobe apps; great for editing and refining images (fill, remove); outputs covered by Adobe license for commercial use	Image gen not as advanced as peers; mainly for Adobe users; beyond small free quota requires CC subscription	Graphic designers & marketers using Photoshop/ Illustrator ; adding AI into design workflows
DreamStudio (Stable Diffusion)	Open-Source (Stable Diffusion)	High variability; requires tuning for best results (can achieve realism or art with the right settings)	Yes - ~100 free credits on signup	Highly customisable (steps, negative prompts, etc.) ; low cost per image; large community & open models	Quality depends on user skill (prompt & settings); default model has significant content moderation ; advanced usage can be technical	AI enthusiasts, artists, or developers wanting control and open-source flexibility

AI Video Generation Tools

These tools deal with **AI-generated video** from text or images, and AI-powered video effects. Text-to-video is an emerging field - these tools create short video clips (often a few seconds) based on prompts.

Sora AI (OpenAI's Text-to-Video Model)



Description: Sora is [OpenAI's](#) experimental text-to-video model, designed to generate short videos (~5-20 seconds) from text descriptions. It can also take image or video inputs to guide the output. Sora is not widely publicly available (in 2025 it's in limited beta for researchers and creators), but it represents cutting-edge quality in AI video. In demos, Sora can handle multiple characters and scene changes in a single clip better than many competitors.

- **Pros: Very high quality output** - among text-to-video tools, Sora's results are some of the most coherent and detailed. It handles **complex scenes** with multiple subjects or actions better than others. Users can specify camera movements in the prompt (e.g. "aerial pan around...") to influence how the scene is shot. It's also more consistent with physics and object permanence than earlier models (though still an area for improvement). Essentially, Sora produces videos that can be surprisingly realistic or on-target for the prompt when it works.
- **Cons:** In usage, some issues include random **scene cuts** in the generated video (the video might suddenly change scene halfway). It still has some content moderation (less strict than DALL-E's, but more than completely open platforms)- certain prompts might be disallowed or result in bland outputs. Additionally, like many generative video models, physics (e.g. realistic motion for running water, or natural human movement) isn't perfect.
- **Pricing:** For significant use, ~[\\$200/month](#), implying a high cost for whatever private beta or API exists. If/when Sora launches publicly, it might come via an OpenAI service (possibly usage-based pricing similar to their image or text models). For now, we can consider that **access is limited to invitees**, and presumably it requires considerable compute credits.
- **Use Cases:** *Future-looking:* generating short **video clips or cinematics** from scratch for creative projects. For instance, storyboarding scenes for films/games, creating ads or visual effects prototypes. In current form, likely used by researchers and studios to explore AI-driven video content. If it becomes widely available, many use cases include quickly visualising an idea in motion (versus a still image), generating background video loops for designs, or video content for social media without filming.
- **Official Info:** [OpenAI Sora](#)

Hailuo AI



Description: [Hailuo](#) (or "Hailuo AI") is a text-to-video model/platform that takes a text prompt and generates a short video (often ~6 seconds). It provides 720p clips at 25 FPS focusing on turning simple prompts into quick videos. It's one of the early consumer-facing AI video generators. Hailuo is notable for a generous free tier, making it a popular choice for users to experiment with AI video.

- **Pros: Lots of free credits** – Hailuo gives new users **1000 credits upfront and 100 credits daily for free**. This means you can generate many videos without paying, which is great for learning and casual use. It follows prompt instructions **better than many competitors**, so the video content tends to match the described scene fairly well. It's also one of the **fastest** generators – around 30 seconds to generate a 6-second video clip, which is quick in this field. So iteration is relatively speedy.
- **Cons:** The video **quality is not top-tier** – outputs "lack polish" and often have artifacts or jitter. While it's improving, you can usually tell it's AI-generated and not a real video. The videos are very short (5 seconds for many prompts, up to 6s), and attempts at longer sequences tend to break. Also, Hailuo's physics and motion can be off (common issues: weird object distortions). It may randomly cut scenes or introduce odd transitions as well. Essentially, it's great for a quick demo, but not production-ready for high quality needs.
- **Pricing: Free plan** available, which includes the daily free credits (100/day). This is enough for a few short videos per day (Hailuo stated ~30 credits per video generation, so ~3 videos a day free). They also have **paid plans** – though specifics aren't in the keynote, the service likely offers subscriptions or top-ups for more credits. According to third-party info, [Hailuo](#) has both free and paid options. We can infer that after using free credits, users can purchase additional credits or subscribe for higher limits (pricing was not explicitly given, but expect tiers for enthusiasts vs. professionals). In summary, one can do a lot for free, and scale up with paid credits if needed.
- **Use Cases:** Great for those who want to **experiment with AI video for fun** – making short clips from imaginative prompts (like "an astronaut dancing on the moon" as a 6s video). Also useful for prototyping video ideas or adding a quick video element to a presentation. Since quality isn't production-grade, it's more for **creative experimentation, education, or social media snippets**. Its ease of use and free credits also make it suitable for students or hobbyists to dip their toes into AI video generation.
- **Official Site:** [hailuoai.video](#) (Hailuo AI's website; also accessible via mobile app).

Runway ML (Gen-2)



Description: [Runway](#) ML is a platform providing various AI tools for creatives. Its **Gen-2** model is a text-to-video system that can create short new videos from a text prompt, or modify existing images/videos. Runway was one of the first to launch text-to-video (Gen-1 and Gen-2). Besides generation, Runway ML offers video editing effects, background removal, and more in a web app.

- **Pros:** Runway's Gen-2 is **one of the first available text-to-video AIs** to the public. The platform is user-friendly: you can input text, get a short video, and also do **video-to-video** (apply a style or transform to an uploaded video). It's a pioneer, so it has a head start in features and has improved with user feedback. Runway also integrates into workflows with its web editor – you can chain AI effects (e.g., generate a video, then use AI to remove background, etc.). Essentially, it's a versatile toolkit.
- **Cons:** Despite being first-mover, the **visual quality is not up to the mark** compared to newer models. Videos often look low fidelity or abstract. The keynote bluntly noted "visual quality is simply not up to the mark". Also, Runway's generous early free tier changed – now there's a credit system. The **cost can be high** for many generations: Gen-2 was reported at ~\$0.05 per 4-second render via their pricing. The subscription tiers ([Standard \\$12/mo](#), [Pro \\$28/mo](#), etc.) include some credits, but for extensive use it can get pricey given the quality trade-off. Another issue: after using included credits, further generations require buying more, which frustrated some users who expected "unlimited". So cost vs. output quality is a concern. Lastly, for now clips are short (a few seconds), similar to others.
- **Pricing: Free plan:** Runway offers a Basic free tier with [125 credits](#) one-time (roughly 25 seconds of video generation), but with watermarks and limited resolution. To do more, the **Standard plan is \$12/user/month** (billed annually) which gives 625 credits/month (~125 seconds of Gen-2 video). The **Pro plan is \$28/month** (annual) for 2250 credits (~450 seconds). On these plans, you can also purchase extra credits or use "slow mode" generation without credits in some cases. Essentially, expect around **\$10-\$30/month** for moderate use. (There was also an unlimited \$100/mo plan with conditions, but it might be throttled as per community discussions.)
- **Use Cases:** As an early tool, Runway Gen-2 is useful for **artistic or abstract video snippets** – for example, generating an arty background video for a design, or testing out generative concepts. It's also popular among **filmmakers for pre-visualization**: you can storyboard ideas by generating rough motion scenes. Additionally, its integrated editor means you can composite the AI video with other media easily, so video editors might use it to generate elements or fill gaps. If absolute realism isn't required, it can be a creative tool for unique, dream-like visuals.
- **Official Site:** [runwayml.com](#) (Gen-2 available in the web app; requires signup).


Kling (Kuaishou's Video Generator)



Description: Kling is a [text-to-video](#) model developed by Kuaishou (a major Chinese social media company). It generates short videos from prompts and also offers an image-to-video capability. Kling is known for striving for higher video quality and has unique features for more precise control (like special effect brushes and lip-sync). It's accessible via a web interface.

- **Pros: High prompt accuracy** - Kling's outputs align closely with the text description, thanks to advanced prompt parsing. It also offers **more user controls** than most: features like a "*motion brush*" to guide movement, or lip-sync to align generated characters with desired audio/dialogue. This level of control is unique and useful for creators who want to fine-tune the result. Kling's quality is considered nearly on par with top models like Sora, but at a lower cost, making it attractive (especially in the Chinese market). Another pro: it's backed by a big company (Kuaishou), indicating active development and integration potential with social media.
- **Cons: Long wait times** - the keynote example noted ~8 minutes to generate the example video, which is significantly slower than some others (likely due to high demand or complex processing). Also, some users observed a "**Chinese look**" in outputs for non-Chinese contexts- possibly a bias in training data causing characters or styles to have an East Asian appearance even if prompt is different. This could require prompt adjustments to avoid. Since it's relatively new to international users, documentation and community help might be limited in English. Lastly, like all these tools, it's not perfect in rendering and can have artifacts.
- **Pricing:** Free users get **~66 credits per day** to generate videos. This is sufficient for one short clip of ~5-10 seconds length per day, since a 5s video might cost ~50 credits (depending on resolution/quality). **Standard plan (~\$10/mo)** provides ~660 credits/month (about 10x the free volume) and removes watermarks and length limits. The **Pro plan (\$37/mo)** gives 3,000 credits and priority processing, while **Premier (\$92/mo)** yields 8,000 credits for heavy user
- **Use Cases:** Content creators who want to generate **short videos or special effects** for social media. Kling's lip-sync and motion control could be used to create AI-driven **characters or avatars that speak lines** - useful for virtual influencers or entertainment. Also, advertisers could try it for quick video ad ideas. Its precise control features make it suitable for those willing to invest time to direct the AI output for more specific results (almost like directing a tiny AI film).
- **Official Site:** <https://klingai.com/>

Pika Labs (Pika AI)



Description: Pika Labs (often just “Pika”) is an AI tool focused on **video-to-video transformations** and special effects. Rather than generating a video wholly from text, Pika lets you upload an existing video or image and apply “Pikaffects” like inflate, melt, explode, or turn things into different materials (e.g., make an object look like cake). It also has text-to-video generation and “Pikadditions” to insert objects into videos. It’s a popular, fun tool for creating viral short videos with eye-catching transformations.

- **Pros: One of the easiest AI VFX tools to use** – the interface is simple: choose an effect, upload media, enter prompt, and it outputs a short video with that effect. It offers very unique and **innovative features** not found elsewhere, like the ability to “cake-ify” someone (make them look like cake) or “puff” an object like a balloon. It’s essentially bringing high-end video effects to casual users via AI. Pika’s effects often look creative and can be quite entertaining, fueling a trend of viral videos. Also, Pika has a free tier, and its **Basic plan is free** with some credits, lowering the barrier to try it.
- **Cons: Inconsistent video quality** – results can vary, sometimes the effect looks convincing, other times glitchy. It’s not at a Hollywood special effects level yet (cannot replace professional VFX for long or high-res videos). Outputs are currently low resolution (720p in many cases) and limited to just a few seconds. Often there are **long wait times** due to popularity and processing load. Another limitation: **very short duration** – many functionalities only let you modify about 5 seconds of video. So you can’t apply these effects to a full minute video in one go. If you need longer, you’d have to stitch multiple segments.
- **Pricing: Free Basic plan:** Pika now offers 80 video credits per month for free users, which allows for a few short video generations depending on the model and features used (e.g., Turbo model costs 5 credits per video, while the 2.1 model with Pikascenes can use up to 60 credits). Paid plans include **Standard at \$8/month** (billed annually) with 700 monthly credits, **Pro at \$28/month** (2,300 monthly credits), and **Fancy at \$76/month** (6,000 monthly credits). All paid plans remove watermarks, unlock HD downloads, and include commercial use rights. Video credit consumption varies based on effects and duration, so actual usage per clip may differ. For most casual users, the Free or Standard plan provides enough to explore and create a handful of videos each month.
- **Use Cases:** Social media content creation – making those **trendy short videos with wild transformations** (like personal TikTok or Instagram clips that stand out). Marketers might use it for attention-grabbing visuals in campaigns. It’s also useful for **creative storytelling**, e.g., taking a still photo and turning it into a dynamic video with an effect (imagine a product photo melting into liquid, etc.). Essentially, Pika shines in **enhancing or remixing existing media** into something fun or surreal. It’s less about accurate rendering and more about artistic, playful effects.
- **Official Site:** pika.art (Web app; also available as mobile app)

Comparison - AI Video Tools:

Tool	Type	Notable Strengths	Limitations	Free Tier	Best Use Case
Sora AI	Text-to-Video (OpenAI research)	Highest quality, complex scenes; can specify camera moves	Limited access (closed beta); very expensive for heavy use; minor issues with cuts/physics	No	Cinematic video generation , high-end creative studios
Hailuo AI	Text-to-Video	Generous free credits; fast generation, good prompt alignment	Short 5-6s videos; average quality (artifacts, not polished)	Yes - 100 credits per day for free	Fun short clips , prototyping video ideas for personal use
Runway Gen-2	Text+Image-to-Video	Easy web interface; integrated editing tools (compositing, etc.); early mover experience	Visual fidelity is low for photo-realism; pay-per-credit can get pricey	Yes - 125 credits one-time free	Artistic or experimental videos ; quick storyboarding with AI visuals
Kling	Text-to-Video	High accuracy to prompts; unique control features (lip-sync, motion brush); nearly Sora-level quality at lower cost	Long waits (minutes) for output; style bias ("Chinese look") in some outputs; availability mostly in China	Yes - 66 credits per day for free	Directed video scenes : when you want more control over AI video content (e.g., dialogue scenes, specific movements)

Pika Labs	Video-to-Video & Effects	Very creative effects and edits (inflate, melt, etc.); extremely easy to use; free tier available	Only 5-10 sec clips; inconsistent output quality; output resolution capped (720p free)	Yes - 80 credits per month for free	Viral social media content , adding quirky effects to images/ videos for marketing or fun
------------------	--------------------------	--	--	--	--

AI Audio Generation Tools

AI isn't just text and images – these tools generate audio, including music and speech.

Suno AI (Music Generation)



Category: AI Music Generation (Text-to-Song)

Description: Suno AI is an AI platform for generating **songs with vocals** from text prompts. It can produce music in various genres complete with instrumental and singing voices. Suno's model (often accessed via the Suno Studio web app) can create songs up to about 4 minutes long based on the description of mood, genre, and lyrics. It's one of the leading tools for AI-generated music with vocals.

- **Pros:** Can produce surprisingly **realistic-sounding songs** (with voices) up to several minutes. It's extremely easy to use – just enter a description of a song (and optionally specify an artist style or provide lyrics) and it generates a sample. Suno provides **generous free usage** for experimentation: 50 credits per day free, which is about 10 songs daily. The AI vocals, while not a specific famous singer, sound coherent as singing, which is impressive. Good for exploring different genres; it tends to capture the vibe (e.g., "a Bollywood love song" prompt yields something that indeed sounds like a Bollywood track)
- **Cons:** The generated vocals have limited variety – currently all songs have a similar sounding singer voice, which users describe as an **"autotuned" effect** across outputs. So while the tune and lyrics change, you may notice a same-voice syndrome. Also, free users cannot use songs commercially (no commercial license on free tier). You'd need a paid plan to legally publish or monetise the music. The AI might produce somewhat generic lyrics or repetitive phrases unless you guide it with specific lyrics. Quality-wise, it's good for basic use, but it won't match a professional studio recording (think of it as a high-quality demo).
- **Pricing:** Free tier: ~50 credits/day, which allows about **10 songs per day for free**. Paid plans: **Pro plan \$10/month** for 2,500 credits/month (~500 songs) and commercial usage rights. There is also a **Premier \$30/mo (or \$24/mo annual)** for 10,000 credits (~2,000 songs). These plans allow using the generated music on platforms like Spotify, YouTube, etc., without copyright issues. In essence, \$10/mo gets most enthusiasts plenty of music generation capacity.
- **Use Cases:** Musicians and creators can use Suno to **generate melodies and song ideas** quickly, potentially overcoming writer's block by seeing what the AI comes up with. Content creators can make background music or theme songs for videos/podcasts without needing a composer (ensuring to get a pro plan if commercial). It's also fun for personal entertainment – for example, you can generate a custom birthday song or a parody by prompting a certain style. Educators might use it to demonstrate music genres or create examples. Essentially, it democratizes making a simple song for those who aren't musically trained.

- **Official Site:** suno.ai (Sign-up for Suno Studio; free and paid options available).

Eleven Labs (Voice Generation & Cloning)



Category: AI Text-to-Speech (Voice Synthesis)

Description: Eleven Labs is a state-of-the-art AI voice generator known for its **ultra-realistic text-to-speech**. It can convert written text into speech that sounds like a human, with options to clone voices or create custom synthetic voices. It supports multiple languages and has become a go-to for high-quality AI narration or voiceover.

- **Pros:** Widely regarded as the **#1 tool in the market for text-to-speech** realisms. The voices it generates (including the default ones) are extremely lifelike, capturing nuances like tone, pacing, and even breathing or laughter in some cases. It's very useful for short, high-quality **voiceovers or narration** – e.g., you can have a script read in a compelling voice that sounds nearly like a real person. Eleven Labs also allows voice cloning: with a sample of a real person's voice (with permission), you can create a custom AI voice that speaks any text, which is how the keynote presenter cloned his own voice in English and Hindi. Multilingual support means you can generate speech in different languages or accents with the same voice.
- **Cons:** Generating and refining voices can be **time-consuming**, especially if you need to get the intonation just right. While it's great for short clips, longer narration might require manual correction (like adjusting where to pause or emphasise). It's also relatively **expensive** for heavy use: the standard plan gives 100k characters for \$22, which is roughly ~1.5 hours of spoken audio file-. For content creators wanting to voice many hours of content, costs add up unless you invest in higher tiers or the new Pro plan. Additionally, **monetising cloned voices is tricky** – using a cloned celebrity voice for content would violate terms; even cloning your own and using it extensively might face platform moderation issues (and as the slide notes, it may be impractical to monetise AI voices on, say, YouTube due to policies or detection). Essentially, it's a fantastic tool but not a magic bullet for unlimited free narration.
- **Pricing:** **Free plan:** allows 10,000 characters per month (~5-6 minutes of audio) for testing. **Starter plan:** \$5/month for 30,000 characters (~15-20 min). The popular **Creator plan:** \$22/month for 100,000 characters (~1.5 hrs). Higher plans: **Pro \$99/mo** for 500k chars, **Scale \$330/mo** for 2M chars. (Prices might have slight variations by region, but \$20 for 100k chars is confirmed) Additional usage beyond included chars is billed per 1k chars (around \$0.30 per 1k on Creator plan). Voice cloning is available from Starter upward (Starter includes instant cloning, Creator adds higher-quality cloning). *Tip:* 100k characters is roughly a 1-hour podcast; so \$22 for an hour of audio is the ballpark.

- **Use Cases: Content creators and publishers** using AI voices for audiobooks, YouTube narration, podcasts (with caution to not violate policies). Also great for prototyping voiceovers in video production before hiring voice actors – you can generate a temp track. Game developers or app developers might use it for NPC dialogue or accessibility features (reading text aloud). Individuals can have documents or articles read to them in a natural voice. And of course, it's popular for fun voice cloning experiments (dubbing yourself in another language or making a favourite character read a script). Essentially, Eleven Labs is for anyone needing high-quality spoken audio from text quickly.
- **Official Site:** elevenlabs.io (Has a web portal for voice generation; API available for developers).

Comparison – AI Audio Tools:

Tool	Output	Key Strengths	Limitations/ Cons	Free Plan	Best For
Suno AI	Full songs (music + vocals)	Generates <i>music with singing</i> , up to 4 min; easy prompt-based song creation; multiple genres	Vocal style is similar across songs (autotune effect); free tier outputs not for commercial use	Yes – ~10 songs/day free	Musicians/ content creators prototyping songs ; quick background music generation
Eleven Labs	Voice speech	<i>Ultra-realistic speech</i> synthesis; supports voice cloning and multi-lingual; suitable for professional voiceovers	Costly for long-form use ; careful use needed for cloned voices (licensing, platform rules)	Yes – 10k chars (~5 min) free	Voiceovers, audiobooks, narration where quality is crucial; creating custom AI voices for characters or accessibility

AI Coding Assistant

Cursor AI (AI Code Editor)



Description: Cursor is an AI-powered code editor (an entire IDE) that helps developers write and refactor code using natural language. It's like having GitHub Copilot or ChatGPT built directly into a coding environment, with a chat interface that has full context of your project. Cursor supports multiple languages and allows you to ask it to generate functions, explain code, fix errors, etc., all within an editor that runs on your machine.

- **Pros: Deep integration with code** – Cursor can take into account your entire codebase when making suggestions or changes, not just a single file. This means it can do things like update all relevant files for a change, or find bugs that span files (a big advantage over simpler code assistants). It's very **fast and intelligent** in autocompleting code or writing new code from scratch based on a comment/instruction. Many users report dramatically increased productivity, as it can scaffold large parts of an application quickly and then you just tweak it. It supports instructing it in plain English, like "refactor this function for efficiency" and it will do so. The UI is similar to VS Code, so familiar to developers.
- **Cons:** As with any AI code tool, it sometimes produces incorrect or suboptimal code – you can't blindly trust every output (especially for critical logic). It might refuse some requests (it has some safeguards, e.g. it won't just write an entire app if it thinks it's too large a request). For very large projects, there might be performance considerations when it's analysing everything. Also, it's a standalone editor – while it's good, some developers have their own highly customised environment and may not switch fully to Cursor. It launched in 2024, so it's relatively new; minor bugs or missing features in the editor part might exist as it's not as mature as, say, VS Code in terms of extensions or debugging features (though it is evolving fast). There's also a privacy consideration: it uses external AI models (like OpenAI's) so code is sent to their servers (important for proprietary code – though they do offer self-hosting for enterprise).
- **Pricing:** As of March 2025, Cursor offers a **Hobby plan that is free for personal use**, providing users with 2,000 completions and 50 slow premium requests per month. For advanced features, the **Pro plan is available at \$20 per month**, offering unlimited completions, 500 fast premium requests per month, and unlimited slow premium requests. The **Business plan, priced at \$40** per user per month, includes all Pro features plus enforced privacy mode organization-wide, centralized team billing, an admin dashboard with usage statistics, and SAML/OIDC SSO. These plans aim to keep Cursor accessible to individual developers, while also catering to enterprise needs through self-hosting options and increased usage limits.
- **Use Cases: Software developers** of all stripes – from students learning to code (Cursor can explain code and suggest fixes, making it a great tutor) to experienced engineers (who can use it

to boilerplate code, write tests, or perform large refactoring swiftly). It's especially useful when working in a new language or framework – the AI can fill in the gaps of your knowledge. Also great for generating documentation or comments for code. Essentially, if you code, an AI pair-programmer like Cursor can speed up tasks (write that function, find that bug) and even serve as a rubber duck to discuss approach. It's like having a smart collaborator who is always available in your IDE.

- **Official Site:** cursor.so

Free Tools List

Many of the above tools offer **free plans or credits**, making them accessible without upfront cost. Here we highlight which AI tools you can use for free (at least to a limited extent) and what you get:

- **[Napkin AI](#)**: *Free during beta* – full functionality currently at no cost. Create up to 3 visuals on the Starter plan even outside beta.
- **[Gamma](#)**: Free plan (no credit card) with 400 AI credits signup bonus. Enough to generate a couple of presentations. Unlimited users and basic exports included.
- **[ChatGPT](#)**: The base ChatGPT models are free to use with some daily limits. These daily limits keep changing from time to time.
- **[Framer](#)**: “Free Forever” tier for websites – allows Framer subdomain and up to 1000 visits/month.
- **[10Web](#)**: Offers a 7-day free trial with full features. Also has a free *Personal* plan option (with very limited resources) for testing.
- **[DALL-E 3](#)**: Free via **[Bing Image Creator](#)** – unlimited usage with a Microsoft account
- **[Leonardo AI](#)**: Comprehensive free use – ~150 tokens every 8 hours (about 15 image generations) perpetually.
- **[Adobe Firefly](#)**: Free plan includes 25 generative credits per month (images have watermark).
- **[DreamStudio \(Stable Diffusion\)](#)**: New users get 25-100 free credits. Enough for dozens of images. After that, pay-per-credit.
- **[Hailuo AI](#)**: Yes – 1000 initial credits + 100 daily free(roughly 3 free videos per day).
- **[Runway ML](#)**: Yes – free tier (Basic) with 125 credits one-time, which equates to ~25 seconds of video gen and some image tools.
- **[Kling](#)**: Yes – 66 free credits daily (roughly three 5 sec videos per day)
- **[Pika Labs](#)**: Yes – free Basic plan with 80 video credits per month (~30 seconds of video effects total).
- **[Suno AI](#)**: Yes – free plan (50 credits daily) giving ~10 songs per day.
- **[Eleven Labs](#)**: Yes – free tier with 10k characters (~5 min speech) per month.
- **[Cursor AI](#)**: Yes – the Cursor code editor is free to download and offers a two week free trail for all its pro features. For the free version, you can have 2000 completions for free.



© 2025 Dhruv Rathee Academy. All rights reserved.

This document and its contents are protected by copyright law and may not be reproduced, distributed, transmitted, displayed, published, or broadcast without the prior written permission of the copyright holder. You may not alter or remove any trademark, copyright, or other notice from copies of the content.